

NUR - Hand-in Exercise 2

Martina Cacciola (s4170814)

March 27, 2024

Abstract

In this document, I present the results for the Hand-in Exercise 2 for the course Numerical Recipes for Astrophysics.

1 Satellite galaxies around a massive central

The exercise is done in the script `satellite.py`. The necessary explanations of the methods used are in the comments of the code. For the purpose of this task, we implement a Random Number Generator (RNG) as follows:

```
1 class CombinedRNG:
2     '''
3     Pseudo-random number generator combining the XOR-shift and Multiply-With-Carry (MWC)
4     methods.
5     We set the seed to the current time in milliseconds, as a source of variability
6     '''
7     def __init__(self, seed = None):
8         # If no seed is provided, use the current time in milliseconds
9         if seed is None:
10             seed = int(time.time() * 1000)
11         # Initialize the states for XOR-shift and MWC methods
12         self.xor_state = seed
13         self.mwc_state = seed
14
15     def xor_shift(self):
16         # XOR-shift method
17         # The seed is converted to a 64-bit unsigned integer
18         # By performing a shift of 21 bits to the left and XORing the result with the
19         # original state
20         self.xor_state ^= (self.xor_state << 21) & 0xFFFFFFFFFFFFFFFF
21         # Then, a shift of 35 bits to the right and again XOR with the previous result
22         self.xor_state ^= (self.xor_state >> 35)
23         # Finally, a shift of 4 bits to the left and XOR with the previous result
24         self.xor_state ^= (self.xor_state << 4) & 0xFFFFFFFFFFFFFFFF
25         return self.xor_state
26
27     def mwc(self):
28         # Multiply-With-Carry (MWC) method
29         # Set the multiplier
30         a = 4294957665
31         # Extract the upper and lower 32 bits of the state
32         mwc_upper = (self.mwc_state & 0xFFFFFFFF00000000) >> 32
33         mwc_lower = self.mwc_state & 0x00000000FFFFFFFF
34         # Calculate the next state by multiplying the lower 32 bits with the multiplier
35         # and adding the upper 32 bits
36         x_next = a * mwc_lower + mwc_upper
37         # Update the state using AND operation with a 64-bit mask
38         # Only the lower 64 bits are kept, the rest are set to zero
39         self.mwc_state = x_next & 0xFFFFFFFFFFFFFFFF
40         return x_next >> 32
41
42     def combined_rng(self):
43         # Combine the XOR-shift and MWC methods by XORing their outputs
44         return (self.xor_shift() ^ self.mwc()) & 0xFFFFFFFFFFFFFFFF
```

```

43
44     def uniform(self, low, high, num_samples):
45         # Generate a list of uniformly distributed random numbers
46         # By scaling the output of the combined RNG to the desired range
47         return [low + (high - low) * self.combined_rng() / 0xFFFFFFFFFFFFFFFF for _ in
48                 range(int(num_samples))]
49
50 # Create an instance of the combined RNG
51 rng = CombinedRNG()

```

satellite.py

For question (a), we do the following:

```

1  ## 1a)
2
3  # Define the parameters
4  a = 2.4
5  b = 0.25
6  c = 1.6
7  xmax = 5
8  Nsat = 100
9
10 def integrand(x):
11     # Take into account singularity at x=0
12     # Consider the volume element is 4*pi*x^2 (4*pi is taken outside)
13     if x == 0:
14         return 0
15     return Nsat * (x/b)**(a-3) * exp(-(x/b)**c) * x**2
16
17 # Implement a numerical integration
18 # Using the trapezoidal rule
19 def trapezoidal_rule(f, a, b, n):
20     """
21     Inputs:
22     f: Function to integrate
23     a, b: Integration limits
24     n: Number of intervals
25     It works by approximating the integral of f(x) between a and b
26     by the sum of the areas of trapezoids formed by the function and the x-axis
27     """
28     # Width of each trapezoid
29     h = (b - a) / n
30     # Initialize the result with the average of the function at the limits
31     result = 0.5 * (f(a) + f(b))
32     for i in range(1, n):
33         # Add the value of the function at each interior point
34         # Each of the function values corresponds to the height of a trapezoid
35         result += f(a + i * h)
36     # Multiply the sum by the width of the trapezoids
37     # This gives the total area under the curve
38     result *= h
39     return result
40
41
42 # Romberg integration
43 def romberg(f, a, b, n):
44     """
45     Inputs:
46     f: Function to integrate
47     a, b: Integration limits
48     n: Number of rows in the Romberg table
49     """
50     R = np.zeros((n, n))
51     # At each iteration, calculates an approximation of the integral
52     # using the trapezoidal rule with 2 ** i intervals
53     # This forms the first column of the Romberg table
54     for i in range(n):
55         R[i, 0] = trapezoidal_rule(f, a, b, 2 ** i)
56     # At each iteration, calculates an improved approximation of the integral
57     # using Richardson extrapolation

```

```

58 # Takes a weighted average of the current approximation and the previous one
59 # The weights are chosen to cancel out as much of the error as possible
60 # This fills the rest of the Romberg table
61 for j in range(1, n):
62     for k in range(j, n):
63         R[k, j] = R[k, j - 1] + (R[k, j - 1] - R[k - 1, j - 1]) / (4 ** j - 1)
64     return R[-1, -1] # Return the last element of the last row of the Romberg table
65
66
67 # Apply the method above to the integral
68 n = 10
69 m = 6
70 R = romberg(integrand, 0, xmax, m)
71 result = R * 4 * pi
72 A = Nsat / result
73 with open("normalization.txt", "w") as file:
74     file.write("The result of the numerical integration is: " + str(A))

```

satellite.py

The normalization factor A we obtain is the following:

```

1 The result of the numerical integration is: 9.05216643171278

```

normalization.txt

. For question (b), we generate 3D satellite positions that statistically follow the satellite profile $n(x)$. We sample them from the probability distribution $p(x)$, using the inverse transform sampling. We do this with the following code:

```

1 ## 1b)
2
3 # Define the number density of galaxies
4 # This represents the number of galaxies per unit volume at a given radius x
5 def n(x):
6     return A * Nsat * (x/b)**(a-3) * exp(-(x/b)**c)
7
8 # Define the number of galaxies N(x) in a shell of radius x and thickness dx
9 # We compute N(x) by multiplying the number density n(x) by the volume element 4*pi*x^2
10 def N(x):
11     return 4 * np.pi * x**2 * n(x)
12
13 # Define the probability distribution function p(x) given that p(x)dx = N(x)dx / Nsat
14 # It gives the probability of finding a galaxy at a given radius x
15 def p_x(x):
16     return N(x) / Nsat
17
18
19 # Implement inverse transform sampling to sample from the distribution
20 def inverse_transform_sampling(pdf, n_samples, x_min, x_max):
21     x_values = np.linspace(x_min, x_max, 10000)
22     cdf_values = np.zeros_like(x_values)
23     # Calculate the cumulative distribution function using numerical integration (
24     # trapezoidal rule)
25     # The cdf at a point x is the integral of the pdf from x_min to x
26     for i in range(1, len(x_values)):
27         cdf_values[i] = cdf_values[i-1] + trapezoidal_rule(pdf, x_values[i-1], x_values[i], 1)
28     cdf_values /= cdf_values[-1] # Normalize to ensure CDF ranges from 0 to 1
29
30     # Generate random numbers uniformly distributed between 0 and 1
31     random_numbers = rng.uniform(0, 1, num_samples=n_samples)
32
33     # For each random number, find the corresponding value of x
34     # such that the CDF of x is equal to the random number
35     # This is done by finding the first x value for which the CDF is greater than the
36     # random number
37     # The corresponding x value is then stored as a sampled point
38     sampled_points = np.zeros(n_samples)

```

```

38     for i in range(n_samples):
39         for j in range(len(cdf_values)):
40             if random_numbers[i] < cdf_values[j]:
41                 sampled_points[i] = x_values[j]
42                 break
43
44     return sampled_points
45
46 # Generate 10,000 sampled points
47 n_samples = 10000
48 x_min, x_max = 10**-4, 5
49 h = (x_max - x_min) / n_samples
50 sampled_points = inverse_transform_sampling(p_x, n_samples, x_min, x_max)
51
52 # Calculate the bin edges
53 bin_edges = np.logspace(np.log10(x_min), np.log10(x_max), 21)
54
55 # Calculate histogram and divide each bin by its width
56 hist, _ = np.histogram(sampled_points, bins=bin_edges, density=False)
57 bin_widths = np.diff(bin_edges)
58 hist_density = hist / (bin_widths * n_samples / Nsat) # Correcting for the
59               normalization offset
60
61 # Plot the analytical function N(x) and the histogram of the sampled points on a log-log
62   scale
63 x_values = np.logspace(-4, np.log10(x_max), 20)
64 N_x_values = [N(x) for x in x_values]
65
66 plt.figure(figsize=(10, 6))
67 plt.loglog(x_values, N_x_values, label='Analytical N(x)', color='r')
68 plt.bar(bin_edges[:-1], hist_density, width=bin_widths, label='Sampled Points Histogram',
69         )
70 plt.xlabel('Relative Radius x')
71 plt.ylabel('Number of Galaxies N(x)')
72 plt.ylim(10**-3, 10**5)
73 plt.legend()
74 plt.title('Analytical N(x) vs Sampled Points Histogram')
75 plt.grid(True)
76 plt.savefig('./plots/my_solution_1b.png', dpi=600)
77
78 ## 1c)
79
80 # Define a selection method following these rules:
81 # Select each galaxy with same probability
82 # Not draw the same galaxy twice
83 # Not reject any drawn galaxy
84 def reservoir_selection(sampled_points, n_galaxies):
85     selected_galaxies = []
86     num_samples = len(sampled_points)
87     for i in range(num_samples):
88         # If we haven't selected n_galaxies yet, add the current galaxy
89         if i < n_galaxies:
90             selected_galaxies.append(sampled_points[i])
91         else:
92             # If we have already selected n_galaxies, decide whether to replace one of
93             them
94             # Generate a random index j between 0 and i
95             j = rng.combined_rng() % (i + 1)
96             # If j is less than n_galaxies, replace the galaxy at index j with the
97             current galaxy
98             if j < n_galaxies:
99                 selected_galaxies[j] = sampled_points[i]
100     return selected_galaxies
101
102 # Define a sorting method (quicksort)
103 # Takes an array and two indices: low (starting index of the array) and high (last index
104   )
105 def quick_sort(arr, low, high):
106     # If the low index is less than the high index, there are elements in the array to

```

```

102     be sorted
103     if low < high:
104         # partition function called to partition the array around a pivot.
105         # The pivot's final position in the sorted array is returned as pi
106         pi = partition(arr, low, high)
107
108         # quick_sort function recursively called for the parts of the array before pi
109         and after pi
110         quick_sort(arr, low, pi - 1)
111         quick_sort(arr, pi + 1, high)
112     return arr
113
114 # This function partitions the array around a pivot
115 def partition(arr, low, high):
116     # Choose the pivot as the median of the first, middle, and last elements
117     mid = low + (high - low) // 2
118     pivot = arr[mid]
119
120     # If the first element is greater than the middle element, swap them
121     if arr[low] > arr[mid]:
122         arr[low], arr[mid] = arr[mid], arr[low]
123
124     # If the middle element is greater than the last element, swap them
125     if arr[mid] > arr[high]:
126         arr[mid], arr[high] = arr[high], arr[mid]
127
128     # If the first element is greater than the middle element, swap them
129     if arr[low] > arr[mid]:
130         arr[low], arr[mid] = arr[mid], arr[low]
131
132     # Now the middle element is the median of the first, middle, and last elements
133     pivot = arr[mid]
134
135     # Swap the pivot with the high element
136     arr[mid], arr[high] = arr[high], arr[mid]
137
138     # i is set to one index less than the low index
139     i = (low - 1)
140     # The array is iterated from the low to the high index
141     for j in range(low, high):
142         # If the element is less than or equal to the pivot, the element is swapped with
143         the one at index i + 1
144         if arr[j] <= pivot:
145             i = i + 1
146             arr[i], arr[j] = arr[j], arr[i]
147
148     # After all elements have been checked, the pivot is swapped with the element at
149     index i + 1
150     # placing the pivot in its correct position in the sorted array
151     arr[i + 1], arr[high] = arr[high], arr[i + 1]

```

satellite.py

We compare the distribution of the sampled points with $N(x)$, the analytical function describing the number of galaxies in a shell of thickness dx at a given distance x . From Fig.1, we can see that the sampled points match the expected distribution. Nevertheless, we can see how the samples are not present in the first interval of x values. Inverse sampling works by transforming uniform random numbers into samples from the desired distribution, using the cumulative distribution function (CDF) of the PDF. The CDF maps values from the sample space to the interval $[0, 1]$, and its inverse can be used to map values from $[0, 1]$ back to the sample space. In our case, we are sampling from $p(x)$, so it is expected for the sampled values to have higher density in regions of higher probability. It is likely to get more samples from the region corresponding to the peak of the distribution, whereas lower values far from the peak have a lower probability of being drawn. This might explain why there is a discrepancy at smaller values in our histogram.

For question (c), we select 100 random galaxies from the ones sampled in point (b), using Reservoir method. It is a sampling algorithm that chooses a random sample, without replacement, of k items from a population of unknown size n in a single pass over the items. The 100 drawn galaxies are ordered from




Figure 1: Plot in logarithmic scale showing $N(x)$, function of the number of galaxies at given distance x , and the histogram of the 10000 sampled points. The sampled points match the analytical distribution.

smallest to largest radius. We do this as follow:

```

1  return (i + 1)
2
3  # Select 100 galaxies
4  n_galaxies = 100
5  selected_galaxies = reservoir_selection(sampled_points, n_galaxies)
6
7  # Sort the galaxies using quicksort from smallest to higher radius
8  sorted_galaxies = quick_sort(selected_galaxies, 0, len(selected_galaxies)-1)
9
10 # Plot the cumulative number of the chosen galaxies
11 fig1c, ax = plt.subplots()
12 ax.plot(sorted_galaxies, np.arange(n_galaxies))
13 ax.set(xscale='log', xlabel='Relative radius',
14        ylabel='Cumulative number of galaxies',
15        xlim=(x_min, x_max), ylim=(0, n_galaxies))
16 plt.savefig('./plots/my_solution-1c.png', dpi=600)
17
18 ## 1d)
19
20 # Define Ridders' method for numerical differentiation
21 def ridders_method(f, x, h, m, tol=1e-10):
22     # Compute first approx with central difference for high h
23     D = [(f(x + h) - f(x - h)) / (2 * h)]
24
25     # Decrease h by a factor of 2 and calculate a new approximation. Repeat until you
26     # have m approximations
27     for i in range(1, m):
28         h /= 2
29         D.append((f(x + h) - f(x - h)) / (2 * h))
30         for j in range(i):
31             D[j] = (4**(j+1) * D[j+1] - D[j]) / (4**(j+1) - 1)
32
33     # Terminate when the improvement over previous best approximation is smaller than
34     # the target error or if error grows
35     best_approximation = None
36     for i in range(1, len(D)):
37         if abs(D[i] - D[i-1]) < tol:

```

```

36         best_approximation = D[i]
37         break
38
39     # Terminate early if the error grows and return the best approximation from
40     before that point.
41     if abs(D[i] - D[i-1]) > abs(D[i-1]):
42         best_approximation = D[i-1]
43         break
44
45     if not best_approximation:
46         print("Tolerance not reached or error grew significantly. Consider increasing m.
47 ")

```

satellite.py

The number of galaxies within a radius x are shown in Fig. 2. The radii selected are starting from higher values because of the previous undersampling at lower values. The increasing trend makes sense as the number of galaxies within a given radius should increase as the radius increases.




Figure 2: Plot showing the cumulative number of galaxies randomly selected with respect to the relative radius x . The number of galaxies progressively increases with larger distances.

For question (d), we compute the derivative of the function $n(x)$ using both the Ridders' method and the analytical procedure. The code is below:

```

1
2     # Now the middle element is the median of the first , middle , and last elements
3     pivot = arr[mid]
4
5     # Swap the pivot with the high element
6     arr[mid], arr[high] = arr[high], arr[mid]
7
8     # i is set to one index less than the low index
9     i = (low - 1)
10    # The array is iterated from the low to the high index

```

```

11     for j in range(low , high):
12         # If the element is less than or equal to the pivot , the element is swapped with
13         the one at index i + 1
14         if arr[j] <= pivot:
15             i = i + 1
16             arr[i], arr[j] = arr[j], arr[i]
17
18     # After all elements have been checked, the pivot is swapped with the element at
19     index i + 1
20     # placing the pivot in its correct position in the sorted array
21     arr[i + 1], arr[high] = arr[high], arr[i + 1]
22     # Return the final position of the pivot
23     return (i + 1)
24
25 # Select 100 galaxies
26 n_galaxies = 100
27 selected_galaxies = reservoir_selection(sampled_points , n_galaxies)
28
29 # Sort the galaxies using quicksort from smallest to higher radius
30 sorted_galaxies = quick_sort(selected_galaxies ,0 ,len(selected_galaxies)-1)
31
32 # Plot the cumulative number of the chosen galaxies
33 fig1c , ax = plt.subplots()
34 ax.plot(sorted_galaxies , np.arange(n_galaxies))
35 ax.set(xscale='log' , xlabel='Relative radius' ,
36        ylabel='Cumulative number of galaxies' ,
37        xlim=(x_min , x_max) , ylim=(0 , n_galaxies))
38 plt.savefig('./plots/my_solution_1c.png' , dpi=600)
39
40 ## 1d)
41
42 # Define Ridders' method for numerical differentiation
43 def ridders_method(f , x , h , m , tol=1e-10):
44     # Compute first approx with central difference for high h
45     D = [(f(x + h) - f(x - h)) / (2 * h)]
46
47     # Decrease h by a factor of 2 and calculate a new approximation. Repeat until you
48     have m approximations

```

satellite.py

The obtained results can be seen in:

```

1 Analytical result: -0.615624912159
2 Numerical result: -0.615624525721

```

derivative.txt

2 Heating and cooling in HII regions

The code relative to this exercise is in the script heating.py. The code relative to point (a) is the following:

```

1 ## 2a)
2
3 # Define the equilibrium function
4 def equilibrium_1(T):
5     term1 = psi * Tc * k
6     term2 = (0.684 - 0.0416 * np.log(T / (1e4 * Z * Z))) * T * k
7     return term1 - term2
8
9 # Measure the time taken and find the root of the equilibrium function using the
10 bisection method
11 def bisection_method(f , a , b , tol=1e-10 , max_iter=100):
12     '''
13     Find the root of the function f(x) = 0 using the bisection method
14     Takes the function f, the endpoints of the initial interval [a, b],
15     and optional arguments tol and max_iter

```



```

16     if f(a) * f(b) >= 0:
17         print("Bisection method fails.")
18         return None, None
19
20     # Initialize values
21     num_steps = 0
22     start_time = time.time()
23     while (b - a) / 2 > tol and num_steps < max_iter:
24         c = (a + b) / 2
25         if f(c) == 0:
26             end_time = time.time()
27             return c, num_steps + 1, end_time - start_time
28
29         # Update the interval
30         if f(c) * f(a) < 0:
31             b = c
32         else:
33             a = c
34         num_steps += 1
35
36     end_time = time.time()
37     return (a + b) / 2, num_steps, end_time - start_time
38
39 # Define the constants
40 k = 1.38e-16 # erg/K
41 aB = 2e-13 # cm^3 / s
42 Z = 0.015 # metallicity
43 Tc = 1e4 * Z**2 # stellar temperature in K
44 psi = 0.929
45
46 # Measure the time taken and find the root of the equilibrium function using the
47 # bisection method
48 a = 1
49 b = 1e7
50 T_eq, num_steps, time_taken = bisection_method(equilibrium_1, a, b)
51
52 if T_eq is not None:
53     output_text = (f"The equilibrium temperature is {T_eq:.2f} K.\n"
54                   f"The bisection method found the root in {num_steps} steps.\n"
55                   f"The time taken was {time_taken:.10f} seconds.\n")
56 else:
57     output_text = "The bisection method did not converge."
58
59 # Writing output to a file
60 with open("2a.txt", "w") as file:
61     file.write(output_text)

```

heating.py

The output of point (a) is in

```

1 The equilibrium temperature is 3.12 K.
2 The bisection method found the root in 56 steps.
3 The time taken was 0.0001101494 seconds.

```

2a.txt

For point (b), we have adopted again the bisection method, which guarantees convergence to a root of the function within a given interval, provided that the function is continuous and changes sign within that interval. It has a linear rate of convergence, which makes it slower with respect to other approaches, e.g. the secant method. We tried also this approach but the time taken for this specific problem was approximately the same and it experienced divergence from the expected root. The code follows the same structure as the one implemented in point (a), now taking into account the density variable n_e . The code is reported below:

```

1 ## 2b)
2
3 # Define the constants

```

```

4 k = 1.38e-16 # erg/K
5 aB = 2e-13 # cm^3 / s
6 Z = 0.015 # metallicity
7 Tc = 1e4 * Z**2 # stellar temperature in K
8 psi = 0.929
9 A = 5e-10 # erg
10 epsilon_CR = 1e-15 # s^-1
11
12 # Define the equilibrium function
13 def equilibrium_2(T, n_e):
14     T4 = T / 1e4
15     term1 = 0.54 * T4**0.37 * aB * n_e * n_e * k * T
16     term2 = A * n_e * epsilon_CR
17     term3 = 8.9e-26 * n_e * T4
18     return term1 - term2 - term3
19
20 # Bisection method
21 def bisection_method_ne(f, a, b, n_e, tol=1e-10, max_iter=100):
22     # Check if the function values at the interval endpoints have the same sign
23     if f(a, n_e) * f(b, n_e) >= 0:
24         print("Bisection method fails.")
25         return None, None, None
26
27     num_steps = 0
28     start_time = time.time()
29     # Iterate until convergence or maximum iterations
30     while (b - a) / 2 > tol and num_steps < max_iter:
31         # Midpoint of the interval
32         c = (a + b) / 2
33         # Check if the midpoint is already the root
34         if f(c, n_e) == 0:
35             end_time = time.time()
36             return c, num_steps + 1, end_time - start_time
37
38         # Check if the root lies between a and c
39         if f(c, n_e) * f(a, n_e) < 0:
40             # Update the upper bound
41             b = c
42         else:
43             # Update the lower bound
44             a = c
45         num_steps += 1
46
47     end_time = time.time()
48     return (a + b) / 2, num_steps, end_time - start_time
49
50 # Densities to consider
51 densities = [1e-4, 1, 1e4]
52
53 with open('2b.txt', 'w') as f:
54     for n_e in densities:
55         # Initial interval for bisection method
56         a_bisection = 1
57         b_bisection = 1e15
58         T_eq_bisection, num_steps_bisection, time_taken_bisection = bisection_method_ne(
59             equilibrium_2, a_bisection, b_bisection, n_e)
60
61         # Write the result to the file
62         if T_eq_bisection is not None:
63             f.write(f"For n_e = {n_e} cm^-3, the equilibrium temperature is {
64                 T_eq_bisection:.2f} K.\n")
65             f.write(f"The bisection method found the root in {num_steps_bisection} steps
66                 .\n")
67             f.write(f"The time taken was {time_taken_bisection:.6f} seconds.\n\n")
68         else:
69             f.write(f"For n_e = {n_e} cm^-3, the bisection method failed to converge.\n\n")

```

heating.py

The output of (b) is in

```
1 For n_e = 0.0001 cm^-3, the equilibrium temperature is 160561340829973.25 K.  
2 The bisection method found the root in 54 steps.  
3 The time taken was 0.000037 seconds.  
4  
5 For n_e = 1 cm^-3, the equilibrium temperature is 34243.13 K.  
6 The bisection method found the root in 83 steps.  
7 The time taken was 0.000057 seconds.  
8  
9 For n_e = 10000.0 cm^-3, the equilibrium temperature is 29.12 K.  
10 The bisection method found the root in 83 steps.  
11 The time taken was 0.000052 seconds.
```

2b.txt

These results represent the equilibrium temperature of the ionized gas in HII regions. This quantity is determined by the balance between heating and cooling processes, which involve various physical phenomena such as free-free emission (Λ_{FF}), heating by cosmic rays (Γ_{CR}), and MHD waves (Γ_{MHD}).

- For **low-density** gas ($n_e = 10^{-4} \text{ cm}^{-3}$), the equilibrium temperature is extremely high (1.60×10^{14} K). This is because, at low densities, the cooling processes (like free-free emission which depends on the square of the density) are inefficient, and the gas remains hot.
- For **intermediate-density** gas ($n_e = 1 \text{ cm}^{-3}$), the equilibrium temperature drops to 34243.13 K. At these densities, the cooling processes become more efficient, and the gas can cool down to lower temperatures.
- For **high-density** gas ($n_e = 10^4 \text{ cm}^{-3}$), the equilibrium temperature is quite low (29.12 K). At high densities, the cooling processes are very efficient, and the gas can cool down to very low temperatures.